

Bringing Order to Chaos: Metagenomics Starting from Raw Reads




DALL-E



Johannes Björk
(UMCG/RUG)

Chrats Melkonian
(WUR & UU; Bioinformatics)


Metagenomics: Overview



Metagenomics: Overview (today's focus)



Metagenomics workflow



DNA extraction



Sequencing




Raw reads

- Exploration
e.g., FASTQC



- Processing



Data preprocessing - Adapters

Adapters are nucleotide sequences placed at either one or both ends of the DNA fragments that are being sequenced

They are composed of 3 sections:

- Sequencer binding site (illumina)
- Multiplexing index (P5-P7)
- Sequencing primer binding site (illumina)

They are necessary for sequencing but should be removed early on in data pre-processing steps




- Index sample A
- Index sample B
- DNA sequence


Data preprocessing - Demultiplexing

Demultiplexing tools: Sabre, iDemux etc..

Generally performed by sequencing companies before sending the data. Good to know what it is to be able to spot it in QC.



Data preprocessing - Adapter trimming



From QC data you may notice that adapters are still present in your sequence. You should remove them either by providing the adapter sequence or using a de-novo search.


Recommended tools:
Trimmomatic, Cutadapt, bbduk, fastp

After adapter removal, rerun QC on the fastq files

Keep an eye out for polyA and polyG sequences


Data preprocessing - Quality filter

Phred quality score – Logarithmic score representing the quality of a nucleotide




Phred Quality Score	Probability of incorrect base call	Base call accuracy
10	1 in 10	90%
20	1 in 100	99%
30	1 in 1000	99.9%

Data preprocessing - Host (& other) removal




Plant reads

e.g. Bowtie2, bbduk,
KneadData, HoCoRT




Metagenomics: Read-based




DNA extraction



Sequencing




Reads




100-150 bp

The read based:
Mapping into
databases



Read-based


Read-based approaches



Align short-reads to different databases containing reference sequences


Read-based ≈ Reference-based

Read-based approaches



Read-based approaches

E.g. HUMAnN,
MetaPathways




Align short-reads to different databases containing reference sequences

Read-based ≈ Reference-based

Read-based approaches


E.g. *ShortBRED*,
ResFinder, *DeepARG*



Align short-reads to different databases containing reference sequences

Read-based ≈ Reference-based

Metagenomics workflow: Assembly-based




DNA extraction



Sequencing



Reads



100-150 bp

Assembly

Contigs



1000+ bp

Assembly-based

The Assembly Problem:

Library of books



Shred all books




Sequencing


Reconstruct each book





Assembly



Repetitive content makes the reconstruction more difficult




Misprints or damaged fragments make reconstruction more difficult



Rare books are difficult or impossible to reconstruct

Back to DNA sequences...



GAAGTAGAATCTTG
GGAAGTAGAATCT
CGGAAGTAGAATCT
TCCGCAAGTAGAATCTC
TCCGGAAGTAGAATCT
ATGTAACACCGGAAGTA
CATGTGTAACTCCGG
TCCATGTGTAACT
CTACCTGTGTAACTC
ACTCCATGTGTAACTCC
AACTCCATGTGT
AAAATCCATATGA
AAAACTCCATGTG
Genome : AAAACTCCATGTGTAACTCCGGAAAGTAGAATCTTG

From these short reads...

reconstruct this

Problem: We don't know where the reads came from in respect to the genome sequence

Genome : AAAACTCCATGTGTAACCTCCGGAAAGTAGAATCTTG

AAAACTCCATGTG
ATGTAACACCGGAAGTA
CATGTGTAACCTCCGG
ACTCCATGTGTAACCTCC
TCCGGAAAGTAGAATCT
CTACCTGTGTAACTC
AAAATCCATATGA
AACTCCATGTGT
GGAAGTAGAATCT
TCCGCAAGTAGAATCTC
TCCATGTGTAACT
CGGAAGTAGAATCT
GAAGTAGAATCTG

From these
short reads...

reconstruct this

Reality: Reads are scrambled ...

AAAACCTCCATGTG
ATGTAACACCGGAAGTA
CATGTGTAACTCCGG
ACTCCATGTGTAACTCC
TCCGGAAGTAGAACATCT
CTACCTGTGTAACTC
AAAATCCATATGA
AACTCCATGTGT
GGAAGTAGAACATCT
TCCGCAAGTAGAACATCTC
TCCATGTGTAACT
CGGAAGTAGAACATCT
GAAGTAGAACATTG

Genome : ?????????????????????????????????? ←.....

From these
short reads...

reconstruct this

Reality: Reads are scrambled AND we don't know the genome sequence

ONE DOES NOT SIMPLY

ASSEMBLE A GENOME




How do we stitch together reads into contigs?

Read A: **AAA**ACTCCATGTG****

Read B: **AAA**ATCCATATGA****


How do we stitch together reads into contigs?

Read A: AAAA**C**TCC**A**T**G**TG



Read B: AAAA**A**TCC**A**T**A**TGA

Suffix - Prefix Overlap



First Law of Assembly

If a suffix of read A is similar to a prefix of read B... then A and B might overlap in the genome

Second Law of Assembly

More coverage leads to more and longer overlaps

More coverage leads to more and longer overlaps

Genome : AAAACTCCATGTGTAACCTCGGAAGTAGAATCTTG
 GGAAGTAGAATCTTG
 TAACTCAGGAAGTAG
 GTGTAACCTCGGA
 TCCATCTGTAACTCC
 AAAACTCCATGTGT

More coverage


Less coverage

More coverage leads to more and longer overlaps

Genome : AAAACTCCATGTGTAACCTCGGAAGTAGAATCTTG
 GGAAGTACAATCTTG
 TAACTCAGGAAGTAG
 GTGTAACCTCCGA
 TCCATCTGTAACTCC
 AAAACTCCATGTGT

 Coverage=3

 Coverage=8



Average coverage:
 $207/35 \approx 6\text{-fold}$

Average coverage:
 $70/35 = 2\text{-fold}$


Overlap graph

Read A: **ACGTA**

Read B: **CGTAC**

Read C: **GTACA**

Contig: **ACGTACA**




Overlap graph

Read A: **ACGTA**

Read B: **CGTAC**


Read C: **GTACA**

Contig: **ACGTACA**




Overlap graph


Read A: **A**CGTA****



Read B: ****CGTAC**C**




Read C: ****CGTACA**C**



Contig: **ACGTACGT**A**** **(With repeats)**


Overlap graph

Read A: **A**CGTA****



Read B: **CGT**A**C**


Read C: **CGT**A**CA**



The introduced repeat "CGT" appears in multiple reads, causing branching paths in the overlap graph


Overlap graph

Read A: **A**CGTA****



Read B: **CGT**A**C**

Read C: **CGT**A**CA**




The introduced repeat "CGT" appears in multiple reads, causing branching paths in the overlap graph

Overlap graph

Read A: **ACGTA**

Read B: **CGTAC**

Read C: **CGTACA**



Path along the
edges with
highest overlap


Not the original contig **ACGTACA**

Overlap graph

Read A: **A**CGTA****

Read B: ****CGTAC**C**

Read C: ****CGTACA****



Alternative path


Original contig **ACGTACGTA**

Overlap graph

Read A: **A**CGTA****

Read B: ****CGTAC**C**

Read C: ****CGTACA****




Not the original contig ****ACGTACA****

Third Law of Assembly

Repeats make assembly difficult


Overlap graph


Original contig **GTACGTACGAT**





Overlap graph


- Each read is a node
- Draw an edge between A and B if suffix of A overlaps with prefix of B
- Contigs are reconstructed by walking along unambiguous paths
- Remove cycles, and at branching paths continue on the edge with the highest overlap


















A **Hamiltonian path** in a graph is a path that visits each **node exactly once**


This path (in this case) is also the **Shortest Common Superstring** which **represents the most compact way to cover all the reads, minimizing redundancy**

Overlap graphs & SCS are not feasible on short-reads

- **Quadratic Complexity in Pairwise Comparisons:** given N reads, this results in $N * (N - 1)$ comparisons, which scales quadratically with the number of reads
- **Finding the Hamiltonian path that gives the exact SCS is NP-hard**
- **Sequencing repeats and errors create ambiguous overlaps**



Repeats introduce ambiguities




Modern short-read assemblers use the de Bruijn graph

Reference genome ACTGAGTACCATGGAC

Sequenced reads ACTGAGTAC
CTGAGTACCAT
GAGTACCATGGAC

k-mers ACTG TACC GGAC
CTGA ACCA
TGAG CCAT
GAGT CATG
AGTA ATGG
GTAC TGGA




Scales linearly instead of quadratically

Modern short-read assemblers use the de Bruijn graph

Reference genome ACTGAGTACCATGGAC

Sequenced reads ACTGAGTAC
CTGAGTACCAT
GAGTACCATGGAC



k-mers ACTG TACC GGAC
CTGA ACCA
TGAG CCAT
GAGT CATG
AGTA ATGG
GTAC TGGA



Eulerian path: Each edge is visited exactly once


Short Read Assembly

(read length < repeat length)




Long Read Assembly

(read length > repeat length)




Short Read Assembly

(read length < repeat length)



Long Read Assembly

(read length > repeat length)




Overlap-Layout-Consensus


Assessing assemblies

Quast


- **Number** of contigs
- **Average/median** contig length
- **Min/Max** contig length
- **N50**: The length of the contigs which covers 50% of genome
- **Read recruitment**: Percentage of all reads mapped back to the assembly
- **Evenness in depth along contig**



Metagenomics workflow: Coverage




Metagenomics workflow: Coverage




DNA extraction



Sequencing



Reads



100-150
bp

Mapping

Taxonomy

Function


Genome


Assembly

Scaffolds

1000+ bp

e.g.
Bowtie2
BWA
SAMtools
...






Depth of coverage (mapping depth)

- Average number of times each nucleotide is covered in the assembly
 - Estimate to the abundance of a sequence in the sample

Breath of coverage (covered length)

- Percentage of bases of a targeted genome that are covered with a certain depth
 - Metagenomic assembly quality – percentage of data included in the assembly
 - Identify chimeric regions

Metagenomics workflow: Taxonomy




DNA extraction



Sequencing



Reads



100-150 bp

Mapping

Taxonomy

Function


Genome

Assembly

Scaffolds


1000+ bp

Stay tuned: Tuesday October 15th





A new tool, RAT, expanding taxonomy assignment on all three levels




+ Function

20 metagenomic classifiers compared:
Simon et.al, 2019
(<https://doi.org/10.1016/j.cell.2019.07.010>)



A new tool, RAT, expanding taxonomy assignment on all three levels




Article | [Open access](#) | Published: 20 April 2024

Integrating taxonomic signals from MAGs and contigs improves read annotation and taxonomic profiling of metagenomes

[Ernestina Hauptfeld](#), [Nikolaos Pappas](#), [Sandra van Iwaarden](#), [Bastien L. Snoek](#), [Andrea Aldas-Vargas](#), [Bas E. Dutilh](#) & [F. A. Bastiaan von Meijenfeldt](#)

Nature Communications 15, Article number: 3373 (2024) | [Cite this article](#)

Metagenomics workflow: Function




DNA extraction



Sequencing



Reads



100-150 bp

Assembly


Scaffolds

Mapping

Taxonom

Function


Genome



Stay tuned: Tuesday October 15th

1000+ bp

Metagenomics workflow: Assembled-based analysis



DNA extraction




Sequencing



Reads



100-150 bp



Assembled-based


Mapping
Taxonom
Function
Genome

Assembly

Scaffolds

1000+ bp

Metagenomics workflow: Binning




DNA extraction



Sequencing



Reads



100-150 bp

Assembly

Mapping


Taxonom

Function

Genome

↓

Scaffolds



1000+ bp


Binning = Separation of genomes from metagenomes

- Who is there and what can every individual do?

PI: How difficult would it be?



PostDoc: It's challenging yet fun, and there are plenty of standardized methods available to help!






PhD candidates & MSc interns

Binning






(MAGs)







Tyson et al. 2004






Albertsen et.al, 2013





Eren et.al, 2015




Automatic tools for binning

Tools:

MetaBAT2
Maxbin2
CONCOCT

Aggregate multiple binning results:
e.i. DASTool


New kids in the playground:
Vamb
SemiBin
MetaDecoder



Overview of the MetaBAT pipeline.

(Semi)Automatic pipelines


Atlas




MetaErg



metaGEM




MetaWRAP



>Standardization of metagenomics is (not) required!

In practice



MAG quality assessment

- Single-copy marker genes


1. Completeness/completion

Marker genes are expected to be present in all bacteria

2. Contamination/redundancy

Single-copy genes are expected to be only once



4,022 closed genomes from NCBI



Golden standard: CheckM

<https://merenlab.org>


Metagenomics workflow: MAG-based analysis



Sequencing



Reads




100-150
bp

Assembly
Scaffolds



1000+ bp 75

MAG-based



- Mapping
- Taxonom
- Function
- Genome

MAG/assembly-based vs. read-based v2

Criteria:	MAG/Assembly-based analysis	Read-based analysis ('mapping')
Comprehensiveness	Low/Medium	Low/Medium/High
Community complexity	Low/Medium	High
Novelty	High	None
Computational burden	High	Low
Genome-resolved metabolism	High	Low
Expert manual supervision	High	Low/Medium
Integration with microbial genomics	High	None

Quince et.al., 2017

MAG/assembly-based vs. read-based v2

Criteria:	MAG/Assembly-based analysis	Read-based analysis ('mapping')
Comprehensiveness	Low/Medium	Low/Medium/High
Community complexity	Low/Medium	High
Novelty	High	None
Computational burden	Medium	Low
Genome-resolved metabolism	High	Low
Expert manual supervision	High	Low/Medium
Integration with microbiomics	High	None

Choose according to your question!

Quince et.al., 2017

Which tools to pick?

Microbiome COSI



HOME ISMB 2024 CAMI COMMUNITY

CAMI2

CAMI II Challenge Information

We proudly announce the beginning of the second round of challenges of the Initiative for the Critical Assessment of Metagenome Interpretation (CAMI) and release of the official challenge data sets!

Over the last two years, we received valuable feedback from the community on important challenges in the field and how to design interesting new data sets and challenges. We incorporated many of your suggestions, thanks again! For you to familiarize with data set types and formats, additional exemplary data sets together with accompanying standards of truth have already been made available over the last months. Two multisample "toy" data sets representing microbial communities from different human body sites and from mouse gut are already provided to allow participants to prepare for the challenges (<https://data.cami-challenge.org/>). These practice data sets are generated from known genomes, and therefore reference-based methods (e.g., using genome databases for their analysis) might perform better here than for real shotgun metagenomic data, where a substantial portion of microbial community members have not been sequenced.

The second CAMI challenge datasets will therefore again include new genomes from taxa (at different evolutionary distances) not found in public databases. Furthermore, a new focus will be on establishing the value of long sequencing reads for microbiome research, with data sets providing both long- and short-read data. Lastly, a clinical pathogen discovery challenge will be offered, mimicking an emergency diagnostic situation in the clinic.

Specifically, the second round of CAMI challenges comprise a metagenome assembly, a genome binning, a taxonomic binning and a taxon across several multi-sample data sets from different environments. This includes a marine data set (ended), a high-strain diversity data set pathogen detection challenge (ended). A new round of challenges on a rhizosphere data set has just started in early 2020!

We are looking forward to receiving your submissions!

The CAMI Team

<https://www.nature.com/articles/s41592-022-01431-4>

Metagenomics is a great tool but...

- Abundance is qualitative
 - Not easy to be quantitative with microbial communities
 - ?Integrate metagenomics/barcoding with qPCR, DNA spiking, flow-cytometry and microscopy?
- We are measuring the DNA content, therefore viable & non viable cells
 - RNA, CFUs (if culturable)
- We investigate potential functionality, not activity
 - Multi-omics: Adding layers of information (RNA, protein, metabolites)
- No clue on spatial organization
 - Microscopy


Metagenomics is a great tool but...

- Abundance is qualitative
 - Not easy to be quantitative with microbial communities
 - ?Integrate metagenomics/barcoding with qPCR, DNA spiking, flow-cytometry and microscopy?
- We are measuring the DNA content, therefore viable & non viable cells
 - RNA, CFUs (if culturable)
- We investigate potential functionality, not activity
 - Multi-omics: Adding layers of information (RNA, protein, metabolites)
- No clue on spatial organization
 - Microscopy
- Toooooooooooooo much data....
 - That's why you are here!

Challenges of getting genomes from metagenomes across environments





- 1) Which of those environments have the highest diversity?
- 2) From which we can get most MAGs?




*High quality:
> 90% completeness &
< 10% contamination

MAG:
metagenome-assembled
genome


Challenges of getting genomes from metagenomes across environments: The more data, the better?




Challenges of getting genomes from metagenomes across environments




Challenges of getting genomes from metagenomes across environments




Challenges of getting genomes from metagenomes across environments




Challenges of getting genomes from metagenomes across environments



Challenges of getting genomes from metagenomes across environments




Key events in winemaking



Challenges of getting genomes from metagenomes across environments



high quality MAGs




Avg shannon index

Hauptfeld, E., et. al. (2022). A metagenomic portrait of the microbial community responsible for two decades of bioremediation of poly-contaminated groundwater. In Water Research (Vol. 221, p. 118767). Elsevier BV. <https://doi.org/10.1016/j.watres.2022.118767>


Challenges of getting genomes from metagenomes across environments



high quality MAGs




Avg shannon index



Hauptfeld, E., et. al. (2022). A metagenomic portrait of the microbial community responsible for two decades of bioremediation of poly-contaminated groundwater. In Water Research (Vol. 221, p. 118767). Elsevier BV. <https://doi.org/10.1016/j.watres.2022.118767>

Practicals




Connect to JupyterHub:

https://bioinformatics.nl/biosb_metagenomics

><https://mdehollander.github.io/biosb-metagenomics/>

By Mattias de Hollander

Practicals



Connect to JupyterHub:

https://bioinformatics.nl/biosb_metagenomics

><https://mdehollander.github.io/biosb-metagenomics/>

By Mattias de Hollander

Feeling adventurous?

Explore the microbiome of a deadly toxic cave

-> **Cave expedition tab**

